

Echo: Tour de Code

The complete function-by-function trace of Echo's execution pipeline.

This document traces EVERY function call involved in processing a user action through the Echo engine. File paths and line numbers are accurate as of 2026-01-18.

Annotated with tour guide commentary — insights, patterns, and observations from a detailed code review.

Tour Guide Notes

Welcome to the Echo Tour de Code! I'll be your guide through this remarkable piece of systems engineering.

What strikes me most about Echo's architecture is its **relentless pursuit of determinism**. Every design decision—from content-addressed identities to 20-pass radix sorts—serves the goal of ensuring that the same inputs always produce the same outputs, regardless of execution timing or parallelism.

As we walk through the pipeline, I'll highlight:

- **Clever patterns** that solve subtle problems elegantly
- **Invariants** that must hold for correctness
- **Performance optimizations** hidden in plain sight
- **Architectural decisions** and their trade-offs

Let's begin our journey from intent to commit!

Table of Contents

1. Intent Ingestion
2. Transaction Lifecycle
3. Rule Matching
4. Scheduler: Drain & Reserve
5. BOAW Parallel Execution

6. Delta Merge & State Finalization
 7. Hash Computation
 8. Commit Orchestration
 9. Complete Call Graph
-

1. Intent Ingestion

Entry Point: `Engine::ingest_intent()` File: `crates/warp-core/src/engine_impl.rs:1216`

📍 Tour Guide Notes

This is where user actions enter the system. Notice how Echo treats intents as *immutable, content-addressed* data from the very first moment. The intent bytes are hashed to create a unique identifier, ensuring that duplicate intents are detected automatically—no coordination required.

1.1 Function Signature

```
pub fn ingest_intent(&mut self, intent_bytes: &[u8]) -> Result<IngestDisposition, EngineError>
```

Returns: - `IngestDisposition::Accepted { intent_id: Hash }` — New intent accepted
- `IngestDisposition::Duplicate { intent_id: Hash }` — Already ingested

1.2 Complete Call Trace

```
Engine::ingest_intent(intent_bytes: &[u8])
```

```
[1] compute_intent_id(intent_bytes) → Hash
FILE: crates/warp-core/src/inbox.rs:205
CODE:
    let mut hasher = blake3::Hasher::new();
    hasher.update(b"intent:");
                                // Domain separation
    hasher.update(intent_bytes);
    hasher.finalize().into()           // → [u8; 32]
```



```
[2] NodeId(intent_id)
Creates strongly-typed NodeId from Hash
```



```
[3] self.state.store_mut(&warp_id) → Option<&mut GraphStore>
FILE: crates/warp-core/src/engine_impl.rs:1221
ERROR: EngineError::UnknownWarp if None
```



```
[4] Extract root_node_id from self.current_root.local_id
```

```
[5] STRUCTURAL NODE CREATION (Idempotent)
    make_node_id("sim") → NodeId
        FILE: crates/warp-core/src/ident.rs:93
        CODE: blake3("node:" || "sim")

    make_node_id("sim/inbox") → NodeId
        FILE: crates/warp-core/src/ident.rs:93
        CODE: blake3("node:" || "sim/inbox")

    make_type_id("sim") → TypeId
        FILE: crates/warp-core/src/ident.rs:85
        CODE: blake3("type:" || "sim")

    make_type_id("sim/inbox") → TypeId
    make_type_id("sim/inbox/event") → TypeId

    store.insert_node(sim_id, NodeRecord { ty: sim_ty })
        FILE: crates/warp-core/src/graph.rs:175
        CODE: self.nodes.insert(id, record)

    store.insert_node(inbox_id, NodeRecord { ty: inbox_ty })

[6] STRUCTURAL EDGE CREATION
    make_edge_id("edge:root/sim") → EdgeId
        FILE: crates/warp-core/src/ident.rs:109
        CODE: blake3("edge:" || "edge:root/sim")

    store.insert_edge(root_id, EdgeRecord { ... })
        FILE: crates/warp-core/src/graph.rs:188
        GraphStore::upsert_edge_record(from, edge)
            FILE: crates/warp-core/src/graph.rs:196
            UPDATES:
                self.edge_index.insert(edge_id, from)
                self.edge_to_index.insert(edge_id, to)
                self.edges_from.entry(from).or_default().push(edge)
                self.edges_to.entry(to).or_default().push(edge_id)

    store.insert_edge(sim_id, EdgeRecord { ... }) [sim → inbox]

[7] DUPLICATE DETECTION
    store.node(&event_id) → Option<&NodeRecord>
        FILE: crates/warp-core/src/graph.rs:87
        CODE: self.nodes.get(id)
        IF Some(_): return Ok(IngestDisposition::Duplicate { intent_id })

[8] EVENT NODE CREATION
```

```

store.insert_node(event_id, NodeRecord { ty: event_ty })
NOTE: event_id = intent_id (content-addressed)

[9] INTENT ATTACHMENT
    AtomPayload::new(type_id, bytes)
        FILE: crates/warp-core/src/attachment.rs:149
        CODE: Self { type_id, bytes: Bytes::copy_from_slice(intent_bytes) }

    store.set_node_attachment(event_id, Some(AttachmentValue::Atom(payload)))
        FILE: crates/warp-core/src/graph.rs:125
        CODE: self.node_attachments.insert(id, v)

[10] PENDING EDGE CREATION (Queue Membership)
    pending_edge_id(&inbox_id, &intent_id) → EdgeId
        FILE: crates/warp-core/src/inbox.rs:212
        CODE: blake3("edge:" || "sim/inbox/pending:" || inbox_id || intent_id)

    store.insert_edge(inbox_id, EdgeRecord {
        id: pending_edge_id,
        from: inbox_id,
        to: event_id,
        ty: make_type_id("edge:pending")
    })
}

[11] return Ok(IngestDisposition::Accepted { intent_id })

```

* Clever Pattern

Domain Separation in Hashing

Notice step [1]: the hasher prefixes with `b"intent:"` before the actual data. This is a cryptographic best practice called *domain separation*—it prevents a hash collision between an intent and, say, a node ID that happens to have the same bytes.

Echo uses this pattern consistently:

- `"intent:"` for intent IDs
- `"node:"` for node IDs
- `"type:"` for type IDs
- `"edge:"` for edge IDs

This ensures that even if two different domain values have the same raw bytes, they'll produce different hashes.

▼ Deep Dive

Why Content-Addressed Event IDs?

In step [8], note that `event_id = intent_id`. This is a profound design choice:

1. **Automatic deduplication:** If the same intent arrives twice, it hashes to the same ID, and step [7] catches it.
2. **Reproducibility:** Given the same intent bytes, any node in a distributed system will compute the same event ID.
3. **Auditability:** You can verify an event's integrity by re-hashing its content.

This is the foundation of Echo's deterministic execution model—events are identified by *what they are*, not *when they arrived*.

1.3 Data Structures Modified

Structure	Field	Change
GraphStore	<code>nodes</code>	+3 entries (sim, inbox, event)
GraphStore	<code>edges_from</code>	+3 edges (root→sim, sim→inbox, inbox→event)
GraphStore	<code>edges_to</code>	+3 reverse entries
GraphStore	<code>edge_index</code>	+3 edge→from mappings
GraphStore	<code>edge_to_index</code>	+3 edge→to mappings
GraphStore	<code>node_attachments</code>	+1 (event → intent payload)

📎 Tour Guide Notes

Notice the **four separate edge indices**: `edges_from`, `edges_to`, `edge_index`, and `edge_to_index`. This redundancy enables O(1) lookups in any direction—find edges from a node, to a node, or look up either endpoint given an edge ID. The space cost is modest (pointers/IDs are small), but the query flexibility is enormous.

2. Transaction Lifecycle

2.1 Begin Transaction

Entry Point: Engine::begin() File: crates/warp-core/src/engine_impl.rs:711-719

```
pub fn begin(&mut self) -> TxId {
    self.tx_counter = self.tx_counter.wrapping_add(1); // Line 713
    if self.tx_counter == 0 {
        self.tx_counter = 1; // Line 715: Zero is reserved
    }
    self.live_txs.insert(self.tx_counter); // Line 717
    TxId::from_raw(self.tx_counter) // Line 718
}
```

△ Watch Out

The Zero Invariant

Line 715 is subtle but critical: `TxId(0)` is reserved as an invalid/sentinel value. Without this check, after 2^{64} transactions (admittedly unlikely!), the counter would wrap to zero and potentially confuse code that uses zero to mean “no transaction.”

This is defensive programming at its finest—the cost is one branch that’s almost never taken, but it eliminates an entire class of potential bugs.

Call Trace:

```
Engine::begin()

self.tx_counter.wrapping_add(1)
Rust std: u64::wrapping_add
Handles u64::MAX → 0 overflow

if self.tx_counter == 0: self.tx_counter = 1
INVARIANT: TxId(0) is reserved as invalid

self.live_txs.insert(self.tx_counter)
TYPE: HashSet<u64>
Registers transaction as active

TxId::from_raw(self.tx_counter)
FILE: crates/warp-core/src/tx.rs:34
CODE: pub const fn from_raw(value: u64) -> Self { Self(value) }
TYPE: #[repr(transparent)] struct TxId(u64)
```

 Tour Guide Notes

The `#[repr(transparent)]` on `TxId` is worth noting—it guarantees that `TxId` has exactly the same memory layout as `u64`. This means zero-cost abstraction: you get type safety (can't accidentally pass a `NodeId` where a `TxId` is expected) with no runtime overhead.

State Changes: - `tx_counter`: $N \rightarrow N+1$ (or 1 if wrapped) - `live_txs`: Insert new counter value

2.2 Abort Transaction

Entry Point: `Engine::abort()` **File:** `crates/warp-core/src/engine_impl.rs:962-968`

```
pub fn abort(&mut self, tx: TxId) {
    self.live_txs.remove(&tx.value());
    self.scheduler.finalize_tx(tx);
    self.bus.clear();
    self.last_materialization.clear();
    self.last_materialization_errors.clear();
}
```

 Tour Guide Notes

Abort is refreshingly simple—just remove the transaction from tracking and clear transient state. No rollback needed because Echo hasn't mutated the graph yet! All graph mutations happen atomically during commit. This is a key architectural decision: the graph is effectively immutable until commit time.

3. Rule Matching

Entry Point: `Engine::apply()` **File:** `crates/warp-core/src/engine_impl.rs:730-737`

 Tour Guide Notes

Now we enter the heart of Echo's reactive model. Rules are matched against graph patterns, and when they match, they're enqueued for execution. The beauty is that matching is *pure*—it reads the graph but doesn't modify it.

3.1 Function Signature

```
pub fn apply(
    &mut self,
```

```

tx: TxId,
rule_name: &str,
scope: &NodeId,
) -> Result<ApplyResult, EngineError>

```

3.2 Complete Call Trace

```

Engine::apply(tx, rule_name, scope)

Engine::apply_in_warp(tx, self.current_root.warp_id, rule_name, scope, &[])
FILE: crates/warp-core/src/engine_impl.rs:754-806

[1] TRANSACTION VALIDATION
    CODE: if tx.value() == 0 || !self.live_txs.contains(&tx.value())
    ERROR: EngineError::UnknownTx

[2] RULE LOOKUP
    self.rules.get(rule_name) -> Option<&RewriteRule>
    TYPE: HashMap<&'static str, RewriteRule>
    ERROR: EngineError::UnknownRule(rule_name.to_owned())

[3] STORE LOOKUP
    self.state.store(&warp_id) -> Option<&GraphStore>
    ERROR: EngineError::UnknownWarp(warp_id)

[4] CREATE GRAPHVIEW
    GraphView::new(store) -> GraphView<'_>
    FILE: crates/warp-core/src/graph_view.rs
    TYPE: Read-only wrapper (Copy, 8 bytes)

[5] CALL MATCHER
    (rule.matcher)(view, scope) -> bool
    TYPE: MatchFn = for<'a> fn(GraphView<'a>, &NodeId) -> bool
    FILE: crates/warp-core/src/rule.rs:16-24
    IF false: return Ok(ApplyResult::NoMatch)

[6] CREATE SCOPE KEY
    let scope_key = NodeKey { warp_id, local_id: *scope }

[7] COMPUTE SCOPE HASH
    scope_hash(&rule.id, &scope_key) -> Hash
    FILE: crates/warp-core/src/engine_impl.rs:1712-1718
    CODE:
        let mut hasher = Hasher::new();
        hasher.update(rule_id);                                // 32 bytes
        hasher.update(scope.warp_id.as_bytes());   // 32 bytes

```

```

    hasher.update(scope.local_id.as_bytes()); // 32 bytes
    hasher.finalize().into()

```

[8] COMPUTE FOOTPRINT

```

    (rule.compute_footprint)(view, scope) → Footprint
    TYPE: FootprintFn = for<'a> fn(GraphView<'a>, &NodeId) -> Footprint
    FILE: crates/warp-core/src/rule.rs:38-46

```

RETURNS:

```

    Footprint {
        n_read: IdSet,           // Nodes read
        n_write: IdSet,          // Nodes written
        e_read: IdSet,           // Edges read
        e_write: IdSet,          // Edges written
        a_read: AttachmentSet,   // Attachments read
        a_write: AttachmentSet,  // Attachments written
        b_in: PortSet,           // Input ports
        b_out: PortSet,          // Output ports
        factor_mask: u64,         // O(1) prefilter
    }

```

[9] AUGMENT FOOTPRINT WITH DESCENT STACK

```

    for key in descent_stack:
        footprint.a_read.insert(*key)
    FILE: crates/warp-core/src/footprint.rs:104-107

```

PURPOSE: Stage B1 law - READs of all descent chain slots

[10] COMPACT RULE ID LOOKUP

```

    self.compact_rule_ids.get(&rule.id) → Option<&CompactRuleId>
    TYPE: HashMap<Hash, CompactRuleId>
    ERROR: EngineError::InternalCorruption

```

[11] ENQUEUE TO SCHEDULER

```

    self.scheduler.enqueue(tx, PendingRewrite { ... })

```

```

DeterministicScheduler::enqueue(tx, rewrite)
FILE: crates/warp-core/src/scheduler.rs:654-659

```

```

RadixScheduler::enqueue(tx, rewrite)
FILE: crates/warp-core/src/scheduler.rs:102-105

```

CODE:

```

let txq = self.pending.entry(tx).or_default();
txq.enqueue(rewrite.scope_hash, rewrite.compact_rule.0, rewrite);

```

```

PendingTx::enqueue(scope_be32, rule_id, payload)
FILE: crates/warp-core/src/scheduler.rs:331-355

```

```
CASE 1: Duplicate (scope_hash, rule_id) - LAST WINS
    index.get(&key) → Some(&i)
    fat[thin[i].handle] = Some(payload) // Overwrite
    thin[i].nonce = next_nonce++      // Refresh nonce

CASE 2: New entry
    fat.push(Some(payload))
    thin.push(RewriteThin { scope_be32, rule_id, nonce, handle })
    index.insert(key, thin.len() - 1)
```

★ Clever Pattern

GraphView: The 8-Byte Read-Only Wrapper

Step [4] creates a `GraphView`—and note it's only **8 bytes** and **Copy!** This is just a pointer to the underlying `GraphStore`, but wrapped in a type that only exposes read methods.

This is Rust's type system doing the heavy lifting: you literally *cannot* mutate the graph through a `GraphView`. The compiler enforces read-only access, enabling safe concurrent reads without any runtime checks.

▼ Deep Dive

The Footprint: Declaring Your Intentions

Step [8] is architecturally critical. Before a rule can execute, it must declare its *footprint*—exactly which nodes, edges, and attachments it will read and write.

This enables:

- **Parallel execution:** Rules with non-overlapping footprints can run concurrently
- **Conflict detection:** Rules with conflicting footprints are serialized
- **Determinism:** The scheduler can order rules without knowing their implementation details

The footprint is computed *before* execution, not discovered during execution. This is a constraint on rule authors, but it's what makes the whole system tractable.

★ Clever Pattern

Last-Wins Deduplication

In step [11], notice the “LAST WINS” semantics. If the same (`scope_hash`, `rule_id`) pair is enqueueued twice, the second one *replaces* the first.

Why? Because enqueueing a rule is idempotent: if you match the same rule at the same scope twice in one transaction, you only want to execute it once. The “last wins” ensures the most recent footprint is used (which matters if the graph changed between matches).

3.3 PendingRewrite Structure

File: crates/warp-core/src/scheduler.rs:68-82

```
pub(crate) struct PendingRewrite {
    pub rule_id: Hash,                      // 32-byte rule identifier
    pub compact_rule: CompactRuleId,          // u32 hot-path handle
    pub scope_hash: Hash,                    // 32-byte ordering key
    pub scope: NodeKey,                     // { warp_id, local_id }
    pub footprint: Footprint,                // Read/write declaration
    pub phase: RewritePhase,                 // State machine: Matched → Reserved → ...
}
```

Tour Guide Notes

Notice the dual identity: `rule_id` (32-byte hash) for correctness, and `compact_rule` (u32) for performance. The hash ensures cryptographic uniqueness; the u32 enables O(1) array indexing. This “have your cake and eat it too” pattern appears throughout Echo.

4. Scheduler: Drain & Reserve

Tour Guide Notes

The scheduler is where Echo’s determinism guarantees are forged. No matter what order rules are enqueued, the scheduler produces a *canonical* execution order. This is perhaps the most technically impressive part of the system.

4.1 Drain Phase (Radix Sort)

Entry Point: `RadixScheduler::drain_for_tx()` File: crates/warp-core/src/scheduler.rs:109-113

```
pub(crate) fn drain_for_tx(&mut self, tx: TxId) -> Vec<PendingRewrite> {
    self.pending
        .remove(&tx)
        .map_or_else(Vec::new, |mut txq| txq.drain_in_order())
}
```

Complete Call Trace:

```
RadixScheduler::drain_for_tx(tx)
```

```
self.pending.remove(&tx) -> Option<PendingTx<PendingRewrite>>
```

```

PendingTx::drain_in_order()
FILE: crates/warp-core/src/scheduler.rs:416-446

DECISION: n <= 1024 (SMALL_SORT_THRESHOLD)?
    YES: sort_unstable_by(cmp_thin)
        Rust std comparison sort

    NO: radix_sort()
        FILE: crates/warp-core/src/scheduler.rs:360-413

radix_sort()

    Initialize scratch buffer: self.scratch.resize(n, default)

    Lazy allocate histogram: self.counts16 = vec![0u32; 65536]

    FOR pass IN 0..20: // 20 PASSES

        SELECT src/dst buffers (ping-pong)
        flip = false: src=thin, dst=scratch
        flip = true: src=scratch, dst=thin

        PHASE 1: COUNT BUCKETS
        FOR r IN src:
            b = bucket16(r, pass)
            counts[b] += 1

        PHASE 2: PREFIX SUMS
        sum = 0
        FOR c IN counts:
            t = *c
            *c = sum
            sum += t

        PHASE 3: STABLE SCATTER
        FOR r IN src:
            b = bucket16(r, pass)
            dst[counts[b]] = r
            counts[b] += 1

        flip = !flip

BUCKET EXTRACTION (bucket16):
FILE: crates/warp-core/src/scheduler.rs:481-498

Pass 0: u16_from_u32_le(r.nonce, 0) // Nonce bytes [0:2]

```

```

Pass 1: u16_from_u32_le(r.nonce, 1)      //Nonce bytes [2:4]
Pass 2: u16_from_u32_le(r.rule_id, 0)    //Rule ID bytes [0:2]
Pass 3: u16_from_u32_le(r.rule_id, 1)    //Rule ID bytes [2:4]
Pass 4: u16_be_from_pair32(scope, 15)   //Scope bytes [30:32]
Pass 5: u16_be_from_pair32(scope, 14)   //Scope bytes [28:30]
...
Pass 19: u16_be_from_pair32(scope, 0)   //Scope bytes [0:2] (MSD)

SORT ORDER: (scope_hash, rule_id, nonce) ascending lexicographic

```

★ Clever Pattern

LSD Radix Sort: O(n) Guaranteed

This is a **Least Significant Digit** radix sort—it processes from the least significant bits to the most significant. After 20 passes (160 bits total), the array is sorted by:

1. `scope_hash` (256 bits, but only 128 bits processed — the top 128 bits provide enough entropy)
2. then `rule_id` (32 bits)
3. then `nonce` (32 bits)

Why radix sort instead of comparison sort?

- **Determinism:** Radix sort is inherently stable and makes no comparisons that could be affected by memory layout
- **O(n) complexity:** With fixed key size, radix sort is linear
- **Cache-friendly:** Sequential memory access in each pass

The 1024-element threshold is a practical optimization: for small arrays, the overhead of radix sort exceeds its benefits, so a comparison sort is used instead.

▼ Deep Dive

Why 20 Passes?

Each pass extracts 16 bits (bucket size 65536). To sort by:

- 128 bits of `scope_hash` = 8 passes
- 32 bits of `rule_id` = 2 passes
- 32 bits of `nonce` = 2 passes

Wait, that's only 12 passes! The extra 8 passes cover more of the `scope_hash` for better distribution. Since LSD radix sort processes from least significant to most significant, passes 4-19 progressively refine the scope ordering.

The nonce is processed first (passes 0-1) because it's the tiebreaker—when `scope_hash` and `rule_id` are equal, the nonce determines order, and we want that to be the finest-grained distinction.

4.2 Reserve Phase (Independence Check)

Entry Point: RadixScheduler::reserve() File: crates/warp-core/src/scheduler.rs:134-143

```
pub(crate) fn reserve(&mut self, tx: TxId, pr: &mut PendingRewrite) -> bool {
    let active = self.active.entry(tx).or_insert_with(ActiveFootprints::new);
    if Self::has_conflict(active, pr) {
        return Self::on_conflict(pr);
    }
    Self::mark_all(active, pr);
    Self::on_reserved(pr)
}
```

Complete Call Trace:

```
RadixScheduler::reserve(tx, pr)

    self.active.entry(tx).or_insert_with(ActiveFootprints::new)
    TYPE: HashMap<TxId, ActiveFootprints>
    ActiveFootprints contains 7 GenSets:
        - nodes_written: GenSet<NodeKey>
        - nodes_read: GenSet<NodeKey>
        - edges_written: GenSet<EdgeKey>
        - edges_read: GenSet<EdgeKey>
        - attachments_written: GenSet<AttachmentKey>
        - attachments_read: GenSet<AttachmentKey>
        - ports: GenSet<PortKey>

    has_conflict(active, pr) -> bool
    FILE: crates/warp-core/src/scheduler.rs:157-236

    FOR node IN pr.footprint.n_write:
        IF active.nodes_written.contains(node): return true // W-W conflict
        IF active.nodes_read.contains(node): return true // W-R conflict

    FOR node IN pr.footprint.n_read:
        IF active.nodes_written.contains(node): return true // R-W conflict
        (R-R is allowed)

    FOR edge IN pr.footprint.e_write:
        IF active.edges_written.contains(edge): return true
        IF active.edges_read.contains(edge): return true

    FOR edge IN pr.footprint.e_read:
        IF active.edges_written.contains(edge): return true

    FOR key IN pr.footprint.a_write:
```

```

IF active.attachments_written.contains(key): return true
IF active.attachments_read.contains(key): return true

FOR key IN pr.footprint.a_read:
    IF active.attachments_written.contains(key): return true

FOR port IN pr.footprint.b_in pr.footprint.b_out:
    IF active.ports.contains(port): return true

IF conflict:
    on_conflict(pr)
        FILE: crates/warp-core/src/scheduler.rs:145-149
        pr.phase = RewritePhase::Aborted
        return false

mark_all(active, pr)
FILE: crates/warp-core/src/scheduler.rs:238-278

FOR node IN pr.footprint.n_write:
    active.nodes_written.mark(NodeKey { warp_id, local_id: node })

FOR node IN pr.footprint.n_read:
    active.nodes_read.mark(NodeKey { ... })

... (similar for edges, attachments, ports)

on_reserved(pr)
FILE: crates/warp-core/src/scheduler.rs:151-155
pr.phase = RewritePhase::Reserved
return true

```

Tour Guide Notes

This is classic **two-phase locking** without the locks! The `has_conflict` function implements the conflict matrix:

	Read	Write
Read	OK	CONFLICT
Write	CONFLICT	CONFLICT

Multiple readers are allowed (R-R is OK), but any write conflicts with both reads and writes of the same resource.

4.3 GenSet: O(1) Conflict Detection

File: crates/warp-core/src/scheduler.rs:509-535

```

pub(crate) struct GenSet<K> {
    gen: u32,                                // Current generation
    seen: FxHashMap<K, u32>,                 // Key → generation when marked
}

impl<K: Hash + Eq + Copy> GenSet<K> {
    #[inline]
    pub fn contains(&self, key: K) -> bool {
        matches!(self.seen.get(&key), Some(&g) if g == self.gen)
    }

    #[inline]
    pub fn mark(&mut self, key: K) {
        self.seen.insert(key, self.gen);
    }
}

```

Key Insight: No clearing needed between transactions. Increment `gen` → all old entries become stale.

★ Clever Pattern

Generation-Based Set: Amortized O(1) Clear

This is one of the most elegant patterns in Echo. Instead of clearing the hash map between transactions ($O(n)$ operation), just increment a generation counter!

An entry is “in the set” only if its stored generation matches the current generation. Old entries with stale generations are effectively invisible. The hash map only grows—it’s never shrunk. But since the same keys tend to be accessed repeatedly (temporal locality), the map stabilizes quickly. The payoff is enormous: clearing the “set” is $O(1)$ instead of $O(n)$.

5. BOAW Parallel Execution

Entry Point: `execute_parallel()` File: `crates/warp-core/src/boaw/exec.rs:61-83`

⌚ Tour Guide Notes

BOAW—“Bag of Asynchronous Work”—is where Echo’s determinism meets parallelism. The key insight: *order of execution doesn’t matter if we sort the outputs*. Rules execute in arbitrary order on worker threads, but their outputs are merged canonically.

5.1 Entry Point

```
pub fn execute_parallel(view: GraphView<'_>, items: &[ExecItem], workers: usize) -> Vec<TickDelta>
    assert!(workers >= 1);
    let capped_workers = workers.min(NUM_SHARDS); // Cap at 256

#[cfg(feature = "parallel-stride-fallback")]
if std::env::var("ECHO_PARALLEL_STRIDE").is_ok() {
    return execute_parallel_stride(view, items, capped_workers);
}

execute_parallel_sharded(view, items, capped_workers) // DEFAULT
}
```

5.2 Complete Call Trace

```
execute_parallel(view, items, workers)

execute_parallel_sharded(view, items, capped_workers)
FILE: crates/warp-core/src/boaw/exec.rs:101-152

IF items.is_empty():
    return (0..workers).map(|_| TickDelta::new()).collect()

partition_into_shards(items.to_vec()) -> Vec<VirtualShard>
FILE: crates/warp-core/src/boaw/shard.rs:109-120

Create 256 empty VirtualShard structures

FOR item IN items:

    shard_of(&item.scope) -> usize
    FILE: crates/warp-core/src/boaw/shard.rs:82-92
    CODE:
        let bytes = scope.as_bytes();
        let first_8: [u8; 8] = [bytes[0..8]];
        let val = u64::from_le_bytes(first_8);
        (val & 255) as usize // SHARD_MASK = 255

    shards[shard_id].items.push(item)

let next_shard = AtomicUsize::new(0)

std::thread::scope(|s| { ... })
FILE: Rust std (scoped threads)
```

```

FOR _ IN 0..workers:

    s.spawn(move || { ... }) // WORKER THREAD

    let mut delta = TickDelta::new()
        FILE: crates/warp-core/src/tick_delta.rs:44-52
        CREATES: { ops: Vec::new(), origins: Vec::new() }

    LOOP: // Work-stealing loop

        shard_id = next_shard.fetch_add(1, Ordering::Relaxed)
            ATOMIC: Returns old value, increments counter
            ORDERING: Relaxed (no synchronization cost)

        IF shard_id >= 256: break

    FOR item IN &shards[shard_id].items:

        let mut scoped = delta.scoped(item.origin)
            FILE: crates/warp-core/src/tick_delta.rs:140-142
            CREATES: ScopedDelta { inner: &mut delta, origin, next_op_ix

                (item.exec)(view, &item.scope, scoped.inner_mut())

        INSIDE EXECUTOR:
            scoped.emit(op)
                FILE: crates/warp-core/src/tick_delta.rs:234-239
                CODE:
                    origin.op_ix = self.next_op_ix;
                    self.next_op_ix += 1;
                    self.inner.emit_with_origin(op, origin);

            TickDelta::emit_with_origin(op, origin)
                FILE: crates/warp-core/src/tick_delta.rs:69-75
                CODE:
                    self.ops.push(op);
                    self.origins.push(origin); // if delta_validate

    COLLECT THREADS:
        handles.into_iter().map(|h| h.join()).collect()
        RETURNS: Vec<TickDelta> (one per worker)

```

★ Clever Pattern

Shard-Based Work Distribution

The sharding scheme is beautifully simple: take the first 8 bytes of the scope's NodeId, mask with 255, and you have your shard.

Why 256 shards?

- **Granularity:** Fine enough that work distributes evenly
- **Overhead:** Coarse enough that per-shard overhead is negligible
- **Determinism:** The shard assignment is deterministic (depends only on NodeId)

The work-stealing loop with `AtomicUsize::fetch_add` is lock-free and cache-friendly—each worker claims shards sequentially, minimizing contention.

▼ Deep Dive

Why `Ordering::Relaxed`?

The atomic counter uses `Relaxed` ordering—the weakest memory ordering. This is safe because:

1. Each shard is processed by exactly one worker (no data races)
2. Workers don't need to see each other's results until after `join()`
3. The `join()` itself provides the necessary synchronization

Using `Relaxed` instead of `SeqCst` avoids memory barriers, which can be expensive on multi-core CPUs.

5.3 ExecItem Structure

File: `crates/warp-core/src/boaw/exec.rs:19-35`

```
#[derive(Clone, Copy)]
pub struct ExecItem {
    pub exec: ExecuteFn,           // fn(GraphView, &NodeId, &mut TickDelta)
    pub scope: NodeId,            // 32-byte node identifier
    pub origin: OpOrigin,         // { intent_id, rule_id, match_ix, op_ix }
}
```

📎 Tour Guide Notes

`ExecItem` is `Clone` + `Copy`—it's just a function pointer plus some IDs. This means workers can own their items without any reference counting or synchronization. The `origin` field enables tracing any operation back to the intent and rule that produced it.

5.4 Thread Safety

Type	Safety	Reason
GraphView	Sync + Send + Clone	Read-only snapshot
ExecItem	Sync + Send + Copy	Function pointer + primitives
TickDelta	Per-worker exclusive	No shared mutation
AtomicUsize	Lock-free	fetch_add with Relaxed ordering

6. Delta Merge & State Finalization

📝 Tour Guide Notes

This is where the magic happens: multiple workers produce independent deltas, and we merge them into a single canonical result. The key invariant: *the merge output depends only on the operations, not on which worker produced them or when.*

6.1 Canonical Merge

Entry Point: `merge_deltas()` **File:** `crates/warp-core/src/boaw/merge.rs:36-75`

```
merge_deltas(deltas: Vec<TickDelta>) → Result<Vec<WarpOp>, MergeConflict>
```

[1] FLATTEN ALL OPS WITH ORIGINS

```
let mut flat: Vec<(WarpOpKey, OpOrigin, WarpOp)> = Vec::new();
FOR d IN deltas:
    let (ops, origins) = d.into_parts_unsorted();
    FOR (op, origin) IN ops.zip(origins):
        flat.push((op.sort_key(), origin, op));
```

[2] CANONICAL SORT

```
flat.sort_by(|a, b| (&a.0, &a.1).cmp(&(&b.0, &b.1)));
ORDER: (WarpOpKey, OpOrigin) lexicographic
```

[3] DEDUPE & CONFLICT DETECTION

```
let mut out = Vec::new();
let mut i = 0;
WHILE i < flat.len():

    GROUP by WarpOpKey
    key = flat[i].0
    start = i
    WHILE i < flat.len() && flat[i].0 == key: i++
```

```

    CHECK if all ops identical
    first = &flat[start].2
    all_same = flat[start+1..i].iter().all(|(_, _, op)| op == first)

    IF all_same:
        out.push(first.clone())           // Accept one copy
    ELSE:
        writers = flat[start..i].iter().map(|(_, o, _)| *o).collect()
        return Err(MergeConflict { writers }) // CONFLICT!

    return Ok(out)

```

★ Clever Pattern

Benevolent Coincidence

The merge allows multiple writers to produce the same operation—this is called a *benevolent coincidence*. If two rules independently decide to create the same edge, that's fine! The merge keeps one copy.

But if they produce *different* operations for the same key (e.g., setting an attachment to different values), that's a `MergeConflict`—a bug in the rule definitions.

This policy allows natural redundancy in rule specifications while catching genuine conflicts.

6.2 WarpOp Sort Key

File: crates/warp-core/src/tick_patch.rs:207-287

```

pub(crate) fn sort_key(&self) -> WarpOpKey {
    match self {
        Self::OpenPortal { .. }      => WarpOpKey { kind: 1, ... },
        Self::UpsertWarpInstance { .. } => WarpOpKey { kind: 2, ... },
        Self::DeleteWarpInstance { .. } => WarpOpKey { kind: 3, ... },
        Self::DeleteEdge { .. }       => WarpOpKey { kind: 4, ... }, // Delete before upsert
        Self::DeleteNode { .. }       => WarpOpKey { kind: 5, ... },
        Self::UpsertNode { .. }       => WarpOpKey { kind: 6, ... },
        Self::UpsertEdge { .. }       => WarpOpKey { kind: 7, ... },
        Self::SetAttachment { .. }    => WarpOpKey { kind: 8, ... }, // Last
    }
}

```

Canonical Order: 1. OpenPortal (creates child instances) 2. UpsertWarpInstance 3. DeleteWarpInstance 4. DeleteEdge (delete before upsert) 5. DeleteNode (delete before upsert) 6. UpsertNode 7. UpsertEdge 8. SetAttachment (after skeleton exists)

▼ Deep Dive

Why This Specific Order?

The operation order is carefully chosen to maintain invariants:

1. **OpenPortal first:** Creates warp instances that other ops may reference
2. **Deletes before upserts:** Ensures we don't accidentally delete something we just created (idempotence)
3. **Nodes before edges:** Edges reference nodes, so nodes must exist first
4. **Attachments last:** Attachments reference nodes/edges, so the skeleton must be complete

This ordering means rules don't need to worry about operation sequencing—emit ops in any order, and the merge will sort them correctly.

6.3 State Mutation Methods

File: crates/warp-core/src/graph.rs

```
GraphStore::insert_node(id, record)
LINE: 175-177
CODE: self.nodes.insert(id, record)
```

```
GraphStore::upsert_edge_record(from, edge)
LINE: 196-261
UPDATES:
- self.edge_index.insert(edge_id, from)
- self.edge_to_index.insert(edge_id, to)
- Remove old edge from previous bucket if exists
- self.edges_from.entry(from).or_default().push(edge)
- self.edges_to.entry(to).or_default().push(edge_id)
```

```
GraphStore::delete_node_cascade(node)
LINE: 277-354
CASCADES:
- Remove from self.nodes
- Remove node attachment
- Remove ALL outbound edges (and their attachments)
- Remove ALL inbound edges (and their attachments)
- Maintain all 4 index maps consistently
```

```
GraphStore::delete_edge_exact(from, edge_id)
LINE: 360-412
VALIDATES: edge is in correct "from" bucket
REMOVES:
- From edges_from bucket
```

- From edge_index
- From edge_to_index
- From edges_to_bucket
- Edge attachment

```
GraphStore::set_node_attachment(id, value)
LINE: 125-134
CODE:
None → self.node_attachments.remove(&id)
Some(v) → self.node_attachments.insert(id, v)

GraphStore::set_edge_attachment(id, value)
LINE: 163-172
Same pattern as node attachments
```

△ Watch Out

Cascade Deletes Are Dangerous

`delete_node_cascade` removes not just the node, but all its edges and attachments. This is correct behavior (dangling edges would violate invariants), but rule authors must be aware: deleting a highly-connected node triggers many index updates.

This is why footprints must declare write access to all edges that might be affected—the cascade happens even if the rule only explicitly deletes the node.

7. Hash Computation

📎 Tour Guide Notes

Hashing is Echo's fingerprint technology. The state root captures *what the graph looks like*; the commit hash captures *how we got here*. Both are computed deterministically using BLAKE3, ensuring that identical states produce identical hashes across all nodes in a distributed system.

7.1 State Root

Entry Point: `compute_state_root()` **File:** `crates/warp-core/src/snapshot.rs:88-209`

```
compute_state_root(state: &WarpState, root: &NodeKey) → Hash
```

[1] BFS REACHABILITY TRAVERSAL

```

Initialize:
reachable_nodes: BTreeset<NodeKey> = { root }
reachable_warps: BTreeset<WarpId> = { root.warp_id }
queue: VecDeque<NodeKey> = [ root ]

WHILE let Some(current) = queue.pop_front():

    store = state.store(&current.warp_id)

    FOR edge IN store.edges_from(&current.local_id):
        to = NodeKey { warp_id: current.warp_id, local_id: edge.to }
        IF reachable_nodes.insert(to): queue.push_back(to)

        IF edge has Descend(child_warp) attachment:
            enqueue_descend(state, child_warp, ...)
                Adds child instance root to queue

        IF current node has Descend(child_warp) attachment:
            enqueue_descend(state, child_warp, ...)
```

[2] HASHING PHASE

```

let mut hasher = Hasher::new() // BLAKE3

HASH ROOT BINDING:
    hasher.update(&root.warp_id.0)      // 32 bytes
    hasher.update(&root.local_id.0)     // 32 bytes

FOR warp_id IN reachable_warps: // BTreeset = sorted order

    HASH INSTANCE HEADER:
        hasher.update(&instance.warp_id.0)      // 32 bytes
        hasher.update(&instance.root_node.0)    // 32 bytes
        hash_attachment_key_opt(&mut hasher, instance.parent.as_ref())

    FOR (node_id, node) IN store.nodes: // BTreeset = sorted
        IF reachable_nodes.contains(&NodeKey { warp_id, local_id: node_id }):
            hasher.update(&node_id.0)           // 32 bytes
            hasher.update(&node_ty.0)          // 32 bytes
            hash_attachment_value_opt(&mut hasher, store.node_attachment(node_id))

    FOR (from, edges) IN store.edges_from: // BTreeset = sorted
        IF from is reachable:
            sorted_edges = edges.filter(reachable).sort_by(|a,b| a.id.cmp(b.id))
            hasher.update(&from.0)             // 32 bytes
            hasher.update(&(sorted_edges.len() as u64).to_le_bytes()) // 8 bytes
```

```

FOR edge IN sorted_edges:
    hasher.update(&edge.id.0)           // 32 bytes
    hasher.update(&edge.ty.0)           // 32 bytes
    hasher.update(&edge.to.0)           // 32 bytes
    hash_attachment_value_opt(&mut hasher, store.edge_attachment(&edge.id))

hasher.finalize().into() // → [u8; 32]

```

★ Clever Pattern

BTreeSet/BTreeMap for Determinism

Notice the use of `BTreeSet` and `BTreeMap` throughout. Unlike `HashSet`/`HashMap`, B-tree collections iterate in *sorted order*. This is essential for deterministic hashing—the hash must be the same regardless of insertion order.

The trade-off: B-tree operations are $O(\log n)$ instead of $O(1)$. But for hashing (which happens once per commit), correctness trumps speed.

▼ Deep Dive

Reachability Pruning

The BFS traversal only hashes *reachable* nodes and edges. This means:

1. Garbage (unreachable nodes) doesn't affect the hash
2. Two states with the same reachable structure have the same hash
3. Deleting a disconnected subgraph doesn't change the hash

This is a subtle but important property for garbage collection—you can safely remove unreachable data without affecting consensus.

7.2 Commit Hash v2

Entry Point: `compute_commit_hash_v2()` File: `crates/warp-core/src/snapshot.rs:244-263`

```

pub(crate) fn compute_commit_hash_v2(
    state_root: &Hash,
    parents: &[Hash],
    patch_digest: &Hash,
    policy_id: u32,
) -> Hash {
    let mut h = Hasher::new();
    h.update(&2u16.to_le_bytes());           // Version tag (2 bytes)
    h.update(&(parents.len() as u64).to_le_bytes()); // Parent count (8 bytes)
    for p in parents {
        h.update(p);                      // Each parent (32 bytes)
    }
    h.update(state_root);                  // Graph hash (32 bytes)
    h.update(patch_digest);              // Ops hash (32 bytes)
}

```

```

    h.update(&policy_id.to_le_bytes());           // Policy (4 bytes)
    h.finalize().into()
}

```

Byte Layout:

Offset	Size	Field
0	2	version_tag (0x02 0x00)
2	8	parent_count (u64 LE)
10	32*N	parents[] (N parent hashes)
10+32N	32	state_root
42+32N	32	patch_digest
74+32N	4	policy_id (u32 LE)

TOTAL: 78 + 32*N bytes → BLAKE3 → 32-byte hash

📎 Tour Guide Notes

The version tag (0x02 0x00) is future-proofing: if the commit hash format ever needs to change, the version lets validators distinguish between formats. The “v2” in the function name indicates this is already the second iteration of the format.

7.3 Patch Digest

Entry Point: `compute_patch_digest_v2()` File: `crates/warp-core/src/tick_patch.rs:755-774`

```

fn compute_patch_digest_v2(
    policy_id: u32,
    rule_pack_id: &ContentHash,
    commit_status: TickCommitStatus,
    in_slots: &[SlotId],
    out_slots: &[SlotId],
    ops: &[WarpOp],
) -> ContentHash {
    let mut h = Hasher::new();
    h.update(&2u16.to_le_bytes());           // Format version
    h.update(&policy_id.to_le_bytes());       // 4 bytes
    h.update(rule_pack_id);                 // 32 bytes
    h.update(&[commit_status.code()]);      // 1 byte
    encode_slots(&mut h, in_slots);
    encode_slots(&mut h, out_slots);
    encode_ops(&mut h, ops);
    h.finalize().into()
}

```

8. Commit Orchestration

Entry Point: `Engine::commit_with_receipt()` File: `crates/warp-core/src/engine_impl.rs:837-954`

Tour Guide Notes

This is the grand finale—where all the pieces come together. The commit orchestrator drains the scheduler, reserves resources, executes rules, merges deltas, computes hashes, and records the transaction. Let's trace through every step.

8.1 Complete Call Trace

```
Engine::commit_with_receipt(tx) → Result<(Snapshot, TickReceipt, WarpTickPatchV1), EngineError>
```

[1] VALIDATE TRANSACTION

```
    IF tx.value() == 0 || !self.live_txs.contains(&tx.value()):  
        return Err(EngineError::UnknownTx)
```

[2] DRAIN CANDIDATES

```
    policy_id = self.policy_id                                // Line 844  
    rule_pack_id = self.compute_rule_pack_id()              // Line 845
```

```
    compute_rule_pack_id()  
    FILE: engine_impl.rs:1675-1688  
    CODE:  
        ids = self.rules.values().map(|r| r.id).collect()  
        ids.sort_unstable(); ids.dedup()  
        hasher.update(&1u16.to_le_bytes()) // version  
        hasher.update(&(ids.len() as u64).to_le_bytes())  
        FOR id IN ids: hasher.update(&id)  
        hasher.finalize().into()
```

```
    drained = self.scheduler.drain_for_tx(tx)                // Line 847  
    plan_digest = compute_plan_digest(&drained)            // Line 848
```

[3] RESERVE (INDEPENDENCE CHECK)

```
    ReserveOutcome { receipt, reserved, in_slots, out_slots }  
    = self.reserve_for_receipt(tx, drained)?                // Line 850-855
```

```
    reserve_for_receipt(tx, drained)  
    FILE: engine_impl.rs:970-1042
```

```
    FOR rewrite IN drained (canonical order):
```

```
        accepted = self.scheduler.reserve(tx, &mut rewrite)
```

```

        IF !accepted:
            blockers = find_blocking_rewrites(reserved, &rewrite)

            receipt_entries.push(TickReceiptEntry { ... })

        IF accepted:
            reserved.push(rewrite)
            extend_slots_from_footprint(&mut in_slots, &mut out_slots, ...)

    return ReserveOutcome { receipt, reserved, in_slots, out_slots }

    rewrites_digest = compute_rewrites_digest(&reserved_rewrites) // Line 858

[4] EXECUTE (PHASE 5 BOAW)
    state_before = self.state.clone() // Line 862
    delta_ops = self.apply_reserved_rewrites(reserved, &state_before)?
        apply_reserved_rewrites(rewrites, state_before)
        FILE: engine_impl.rs:1044-1105

        let mut delta = TickDelta::new()

        FOR rewrite IN rewrites:
            executor = self.rule_by_compact(rewrite.compact_rule).executor
            view = GraphView::new(self.state.store(&rewrite.scope.warp_id))
            (executor)(view, &rewrite.scope.local_id, &mut delta)

        let ops = delta.finalize() // Canonical sort

        patch = WarpTickPatchV1::new(policy_id, rule_pack_id, ..., ops)
        patch.apply_to_state(&mut self.state)?
            [delta_validate]: assert_delta_matches_diff(&ops, &diff_ops)

[5] MATERIALIZE
    mat_report = self.bus.finalize() // Line 884
    self.last_materialization = mat_report.channels
    self.last_materialization_errors = mat_report.errors

[6] COMPUTE DELTA PATCH
    ops = diff_state(&state_before, &self.state) // Line 889

    diff_state(before, after)
    FILE: tick_patch.rs:979-1069
    - Canonicalize portal authoring (OpenPortal)

```

```

        - Diff instances (delete/upsert)
        - Diff nodes, edges, attachments
        - Sort by WarpOp::sort_key()

patch = WarpTickPatchV1::new(policy_id, rule_pack_id, ..., ops)
patch_digest = patch.digest()                                // Line 898

[7] COMPUTE STATE ROOT
    state_root = compute_state_root(&self.state, &self.current_root) // Line 900

[8] GET PARENTS
    parents = self.last_snapshot.as_ref().map(|s| vec![s.hash]).unwrap_or_default()

[9] COMPUTE DECISION DIGEST
    decision_digest = receipt.digest()                         // Line 929

[10] COMPUTE COMMIT HASH
    hash = compute_commit_hash_v2(&state_root, &parents, &patch_digest, policy_id)

[11] BUILD SNAPSHOT
    snapshot = Snapshot {
        root: self.current_root,
        hash,                      // commit_id v2
        parents,
        plan_digest,              // Diagnostic
        decision_digest,           // Diagnostic
        rewrites_digest,           // Diagnostic
        patch_digest,              // COMMITTED
        policy_id,                 // COMMITTED
        tx,
    }

[12] RECORD TO HISTORY
    self.last_snapshot = Some(snapshot.clone())             // Line 947
    self.tick_history.push((snapshot, receipt, patch))     // Line 948-949
    self.live_txs.remove(&tx.value())                     // Line 951
    self.scheduler.finalize_tx(tx)                         // Line 952

[13] RETURN
    Ok((snapshot, receipt, patch))

```

* Clever Pattern

State Snapshot Before Mutation

In step [4], notice `state_before = self.state.clone()`. This clone

happens *before* any mutations. Why?

1. Enables `diff_state()` to compute exactly what changed
 2. Supports rollback if execution fails (though this isn't shown)
 3. Provides validation: the delta from execution should match the diff
- The clone is relatively cheap because it's copy-on-write under the hood—most data is shared until mutation.

▼ Deep Dive

Diagnostic vs. Committed Digests

The snapshot contains multiple digests, but only some are “committed” (affect the hash):

- **Committed:** `state_root`, `patch_digest`, `policy_id`, `parents`
- **Diagnostic:** `plan_digest`, `decision_digest`, `rewrites_digest`

Diagnostic digests are for debugging and auditing—they help trace what happened, but don't affect consensus. This separation keeps the consensus-critical path minimal while providing rich observability.

8.2 Commit Hash Inputs

Input	Committed?	Purpose
<code>state_root</code>		What the graph looks like
<code>patch_digest</code>		How we got here (ops)
<code>parents</code>		Chain continuity
<code>policy_id</code>		Aion policy version
<code>plan_digest</code>		Diagnostic only
<code>decision_digest</code>		Diagnostic only
<code>rewrites_digest</code>		Diagnostic only

9. Complete Call Graph

9.1 Full Journey: Intent → Commit

USER ACTION

```
Engine::ingest_intent(intent_bytes)
    compute_intent_id()                                // BLAKE3 content hash
    make_node_id(), make_type_id()                     // Structural IDs
    store.insert_node()                               // Create event node
    store.set_node_attachment()                      // Attach intent payload
```



```

        delta.finalize()                      // Sort ops
        patch.apply_to_state(&mut self.state)

[ MATERIALIZE ]
    bus.finalize()

[ DELTA PATCH ]
    diff_state(&state_before, &self.state)
        Sort by WarpOp::sort_key()
    WarpTickPatchV1::new(...)
        compute_patch_digest_v2()

[ HASHES ]
    compute_state_root(&self.state, &self.current_root)
        BFS reachability
        BLAKE3 over canonical encoding
    compute_commit_hash_v2(state_root, parents, patch_digest, policy_id)
        BLAKE3(version || parents || state_root || patch_digest || policy_id)

[ SNAPSHOT ]
    Snapshot { root, hash, parents, digests..., policy_id, tx }

[ RECORD ]
    tick_history.push((snapshot, receipt, patch))
    live_txs.remove(&tx.value())
    scheduler.finalize_tx(tx)

```

RETURN: (Snapshot, TickReceipt, WarpTickPatchV1)

Tour Guide Notes

And there you have it—the complete journey from user action to committed state. Every step is deterministic, every hash is content-addressed, and the system can be replayed or verified by any node with the same inputs.

The elegance lies in the separation of concerns:

- **Ingestion** is pure data capture
- **Matching** is pure pattern recognition
- **Scheduling** is pure ordering
- **Execution** is pure computation (no side effects escape)
- **Merging** is pure deduplication
- **Hashing** is pure fingerprinting

Each phase can be reasoned about independently, tested independently, and optimized independently. This is the hallmark of well-architected systems.

9.2 File Index

Component	Primary File	Key Lines
Intent Ingestion	<code>engine_impl.rs</code>	1216-1281
Identity Hashing	<code>ident.rs</code>	85-109
Transaction Begin	<code>engine_impl.rs</code>	711-719
Rule Apply	<code>engine_impl.rs</code>	730-806
Footprint	<code>footprint.rs</code>	131-152
Scheduler Enqueue	<code>scheduler.rs</code>	102-105, 331-355
Radix Sort	<code>scheduler.rs</code>	360-413, 481-498
Reserve/Conflict	<code>scheduler.rs</code>	134-278
GenSet	<code>scheduler.rs</code>	509-535
BOAW Execute	<code>boaw/exec.rs</code>	61-152
Shard Routing	<code>boaw/shard.rs</code>	82-120
Delta Merge	<code>boaw/merge.rs</code>	36-75
TickDelta	<code>tick_delta.rs</code>	38-172
WarpOp Sort Key	<code>tick_patch.rs</code>	207-287
State Mutations	<code>graph.rs</code>	175-412
Patch Apply	<code>tick_patch.rs</code>	434-561
Diff State	<code>tick_patch.rs</code>	979-1069
State Root Hash	<code>snapshot.rs</code>	88-209
Commit Hash v2	<code>snapshot.rs</code>	244-263
Patch Digest	<code>tick_patch.rs</code>	755-774
Commit Orchestrator	<code>engine_impl.rs</code>	837-954

Appendix A: Complexity Summary

Operation	Complexity	Notes
<code>ingest_intent</code>	$O(1)$	Fixed structural insertions
<code>begin</code>	$O(1)$	Counter increment + set insert
<code>apply</code>	$O(m)$	$m = \text{footprint size}$
<code>drain_for_tx</code> (radix)	$O(n)$	$n = \text{candidates, 20 passes}$
reserve per rewrite	$O(m)$	$m = \text{footprint size, } O(1) \text{ per check}$
<code>execute_parallel</code>	$O(n/w)$	$n = \text{items, } w = \text{workers}$
<code>merge_deltas</code>	$O(k \log k)$	$k = \text{total ops (sort + dedup)}$
<code>compute_state_root</code>	$O(V + E)$	$V = \text{nodes, } E = \text{edges}$
<code>compute_commit_hash_v2</code>	$O(P)$	$P = \text{parents}$

📎 Tour Guide Notes

Notice that all operations are either $O(1)$, $O(n)$, or $O(n \log n)$ —there’s nothing quadratic or exponential lurking here. The system scales linearly with the amount of work, which is essential for predictable performance. The one potential bottleneck is `compute_state_root` at $O(V + E)$, which traverses the entire reachable graph. For very large graphs, this could become expensive. In practice, graphs are partitioned across warp instances, keeping each traversal manageable.

Appendix B: Determinism Boundaries

Guaranteed Deterministic

- Radix sort ordering (20-pass LSD)
- BTreeMap/BTreeSet iteration
- BLAKE3 hashing
- GenSet conflict detection
- Canonical merge deduplication

Intentionally Non-Deterministic (Handled by Merge)

- Worker execution order in BOAW
- Shard claim order (atomic counter)

▼ Deep Dive

The Determinism Contract

Echo’s determinism guarantee is: *given the same inputs (intents, rules, initial state), the output (commit hash) is identical across all executions.*

This holds even though:

- Workers execute in arbitrary order
- Shards are claimed non-deterministically
- Thread scheduling varies between runs

The canonical merge absorbs this non-determinism, producing a deterministic output from non-deterministic intermediate results. It’s a beautiful example of “eventual determinism”—chaos in the middle, order at the end.

Protocol Constants (Frozen)

- `NUM_SHARDS` = 256
- `SHARD_MASK` = 255
- Shard routing: `LE_u64(node_id[0..8]) & 255`

- Commit hash v2 version tag: 0x02 0x00

△ Watch Out

Protocol Constants Are Sacred

These constants are “frozen”—changing them would break compatibility with existing commits. If you’re tempted to tweak NUM_SHARDS or the shard routing formula, remember: every historical commit was created with these values, and changing them would make replay impossible. Protocol evolution happens through version tags (like the 0x02 in commit hash v2), not by modifying existing constants.

✎ Tour Guide Notes

End of Tour

Thank you for joining me on this journey through Echo’s internals! We’ve seen:

- **Content-addressed everything:** From intents to commits, identity comes from content
- **Deterministic scheduling:** Radix sort + footprints = predictable execution
- **Safe parallelism:** Sharded execution + canonical merge = speed without chaos
- **Cryptographic integrity:** BLAKE3 hashes throughout = verifiable state

Echo is a remarkable piece of engineering—complex enough to solve hard problems, yet built from simple, composable primitives. The code rewards careful study, and I hope these annotations help illuminate the “why” behind the “what.”

Happy hacking!

Document generated 2026-01-18. File paths and line numbers accurate as of this date. Commentary added by your friendly AI tour guide.

