

Echo: Tour de Code

The complete function-by-function trace of Echo's execution pipeline.

This document traces EVERY function call involved in processing a user action through the Echo engine. File paths and line numbers are accurate as of 2026-01-18.

Table of Contents

1. Intent Ingestion
 2. Transaction Lifecycle
 3. Rule Matching
 4. Scheduler: Drain & Reserve
 5. BOAW Parallel Execution
 6. Delta Merge & State Finalization
 7. Hash Computation
 8. Commit Orchestration
 9. Complete Call Graph
-

1. Intent Ingestion

Entry Point: Engine::ingest_intent() **File:** crates/warp-core/src/engine_impl.rs:1216

1.1 Function Signature

```
pub fn ingest_intent(&mut self, intent_bytes: &[u8]) -> Result<IngestDisposition, EngineError>
```

Returns: - IngestDisposition::Accepted { intent_id: Hash } — New intent accepted
- IngestDisposition::Duplicate { intent_id: Hash } — Already ingested

1.2 Complete Call Trace

```
Engine::ingest_intent(intent_bytes: &[u8])

[1] compute_intent_id(intent_bytes) → Hash
    FILE: crates/warp-core/src/inbox.rs:205
    CODE:
        let mut hasher = blake3::Hasher::new();
        hasher.update(b"intent:");           // Domain separation
        hasher.update(intent_bytes);
        hasher.finalize().into()           // → [u8; 32]

[2] NodeId(intent_id)
    Creates strongly-typed NodeId from Hash

[3] self.state.store_mut(&warp_id) → Option<&mut GraphStore>
    FILE: crates/warp-core/src/engine_impl.rs:1221
    ERROR: EngineError::UnknownWarp if None

[4] Extract root_node_id from self.current_root.local_id

[5] STRUCTURAL NODE CREATION (Idempotent)
    make_node_id("sim") → NodeId
        FILE: crates/warp-core/src/ident.rs:93
        CODE: blake3("node:" || "sim")

    make_node_id("sim/inbox") → NodeId
        CODE: blake3("node:" || "sim/inbox")

    make_type_id("sim") → TypeId
        FILE: crates/warp-core/src/ident.rs:85
        CODE: blake3("type:" || "sim")

    make_type_id("sim/inbox") → TypeId
    make_type_id("sim/inbox/event") → TypeId

    store.insert_node(sim_id, NodeRecord { ty: sim_ty })
        FILE: crates/warp-core/src/graph.rs:175
        CODE: self.nodes.insert(id, record)

    store.insert_node(inbox_id, NodeRecord { ty: inbox_ty })

[6] STRUCTURAL EDGE CREATION
    make_edge_id("edge:root/sim") → EdgeId
        FILE: crates/warp-core/src/ident.rs:109
        CODE: blake3("edge:" || "edge:root/sim")
```

```

store.insert_edge(root_id, EdgeRecord { ... })
FILE: crates/warp-core/src/graph.rs:188
GraphStore::upsert_edge_record(from, edge)
FILE: crates/warp-core/src/graph.rs:196
UPDATES:
    self.edge_index.insert(edge_id, from)
    self.edge_to_index.insert(edge_id, to)
    self.edges_from.entry(from).or_default().push(edge)
    self.edges_to.entry(to).or_default().push(edge_id)

store.insert_edge(sim_id, EdgeRecord { ... }) [sim → inbox]

[7] DUPLICATE DETECTION
store.node(&event_id) → Option<&NodeRecord>
FILE: crates/warp-core/src/graph.rs:87
CODE: self.nodes.get(id)
IF Some(_): return Ok(IngestDisposition::Duplicate { intent_id })

[8] EVENT NODE CREATION
store.insert_node(event_id, NodeRecord { ty: event_ty })
NOTE: event_id = intent_id (content-addressed)

[9] INTENT ATTACHMENT
AtomPayload::new(type_id, bytes)
FILE: crates/warp-core/src/attachment.rs:149
CODE: Self { type_id, bytes: Bytes::copy_from_slice(intent_bytes) }

store.set_node_attachment(event_id, Some(AttachmentValue::Atom(payload)))
FILE: crates/warp-core/src/graph.rs:125
CODE: self.node_attachments.insert(id, v)

[10] PENDING EDGE CREATION (Queue Membership)
pending_edge_id(&inbox_id, &intent_id) → EdgeId
FILE: crates/warp-core/src/inbox.rs:212
CODE: blake3("edge:" || "sim/inbox/pending:" || inbox_id || intent_id)

store.insert_edge(inbox_id, EdgeRecord {
    id: pending_edge_id,
    from: inbox_id,
    to: event_id,
    ty: make_type_id("edge:pending")
})

[11] return Ok(IngestDisposition::Accepted { intent_id })

```

1.3 Data Structures Modified

Structure	Field	Change
GraphStore	nodes	+3 entries (sim, inbox, event)
GraphStore	edges_from	+3 edges (root→sim, sim→inbox, inbox→event)
GraphStore	edges_to	+3 reverse entries
GraphStore	edge_index	+3 edge→from mappings
GraphStore	edge_to_index	+3 edge→to mappings
GraphStore	node_attachments	+1 (event → intent payload)

2. Transaction Lifecycle

2.1 Begin Transaction

Entry Point: Engine::begin() File: crates/warp-core/src/engine_impl.rs:711-719

```
pub fn begin(&mut self) -> TxId {
    self.tx_counter = self.tx_counter.wrapping_add(1); // Line 713
    if self.tx_counter == 0 {
        self.tx_counter = 1; // Line 715: Zero is reserved
    }
    self.live_txs.insert(self.tx_counter); // Line 717
    TxId::from_raw(self.tx_counter) // Line 718
}
```

Call Trace:

```
Engine::begin()

self.tx_counter.wrapping_add(1)
Rust std: u64::wrapping_add
Handles u64::MAX → 0 overflow

if self.tx_counter == 0: self.tx_counter = 1
INVARIANT: TxId(0) is reserved as invalid

self.live_txs.insert(self.tx_counter)
TYPE: HashSet<u64>
Registers transaction as active
```

```
TxId::from_raw(self.tx_counter)
FILE: crates/warp-core/src/tx.rs:34
CODE: pub const fn from_raw(value: u64) -> Self { Self(value) }
TYPE: #[repr(transparent)] struct TxId(u64)
```

State Changes: - tx_counter: $N \rightarrow N+1$ (or 1 if wrapped) - live_txs: Insert new counter value

2.2 Abort Transaction

Entry Point: Engine::abort() **File:** crates/warp-core/src/engine_impl.rs:962-968

```
pub fn abort(&mut self, tx: TxId) {
    self.live_txs.remove(&tx.value());
    self.scheduler.finalize_tx(tx);
    self.bus.clear();
    self.last_materialization.clear();
    self.last_materialization_errors.clear();
}
```

3. Rule Matching

Entry Point: Engine::apply() **File:** crates/warp-core/src/engine_impl.rs:730-737

3.1 Function Signature

```
pub fn apply(
    &mut self,
    tx: TxId,
    rule_name: &str,
    scope: &NodeId,
) -> Result<ApplyResult, EngineError>
```

3.2 Complete Call Trace

Engine::apply(tx, rule_name, scope)

```
Engine::apply_in_warp(tx, self.current_root.warp_id, rule_name, scope, &[])
FILE: crates/warp-core/src/engine_impl.rs:754-806
```

```
[1] TRANSACTION VALIDATION
CODE: if tx.value() == 0 || !self.live_txs.contains(&tx.value())
ERROR: EngineError::UnknownTx
```

```
[2] RULE LOOKUP
    self.rules.get(rule_name) → Option<&RewriteRule>
    TYPE: HashMap<&'static str, RewriteRule>
    ERROR: EngineError::UnknownRule(rule_name.to_owned())

[3] STORE LOOKUP
    self.state.store(&warp_id) → Option<&GraphStore>
    ERROR: EngineError::UnknownWarp(warp_id)

[4] CREATE GRAPHVIEW
    GraphView::new(store) → GraphView<'_>
    FILE: crates/warp-core/src/graph_view.rs
    TYPE: Read-only wrapper (Copy, 8 bytes)

[5] CALL MATCHER
    (rule.matcher)(view, scope) → bool
    TYPE: MatchFn = for<'a> fn(GraphView<'a>, &NodeId) -> bool
    FILE: crates/warp-core/src/rule.rs:16-24
    IF false: return Ok(ApplyResult::NoMatch)

[6] CREATE SCOPE KEY
    let scope_key = NodeKey { warp_id, local_id: *scope }

[7] COMPUTE SCOPE HASH
    scope_hash(&rule.id, &scope_key) → Hash
    FILE: crates/warp-core/src/engine_impl.rs:1712-1718
    CODE:
        let mut hasher = Hasher::new();
        hasher.update(rule_id); // 32 bytes
        hasher.update(scope.warp_id.as_bytes()); // 32 bytes
        hasher.update(scope.local_id.as_bytes()); // 32 bytes
        hasher.finalize().into()

[8] COMPUTE FOOTPRINT
    (rule.compute_footprint)(view, scope) → Footprint
    TYPE: FootprintFn = for<'a> fn(GraphView<'a>, &NodeId) -> Footprint
    FILE: crates/warp-core/src/rule.rs:38-46
    RETURNS:
        Footprint {
            n_read: IdSet, // Nodes read
            n_write: IdSet, // Nodes written
            e_read: IdSet, // Edges read
            e_write: IdSet, // Edges written
            a_read: AttachmentSet, // Attachments read
            a_write: AttachmentSet, // Attachments written
            b_in: PortSet, // Input ports
        }
    
```

```

        b_out: PortSet,           // Output ports
        factor_mask: u64,         // O(1) prefilter
    }

[9] AUGMENT FOOTPRINT WITH DESCENT STACK
    for key in descent_stack:
        footprint.a_read.insert(*key)
FILE: crates/warp-core/src/footprint.rs:104-107
PURPOSE: Stage B1 law - READs of all descent chain slots

[10] COMPACT RULE ID LOOKUP
    self.compact_rule_ids.get(&rule.id) → Option<&CompactRuleId>
TYPE: HashMap<Hash, CompactRuleId>
ERROR: EngineError::InternalCorruption

[11] ENQUEUE TO SCHEDULER
    self.scheduler.enqueue(tx, PendingRewrite { ... })

DeterministicScheduler::enqueue(tx, rewrite)
FILE: crates/warp-core/src/scheduler.rs:654-659

    RadixScheduler::enqueue(tx, rewrite)
    FILE: crates/warp-core/src/scheduler.rs:102-105
    CODE:
        let txq = self.pending.entry(tx).or_default();
        txq.enqueue(rewrite.scope_hash, rewrite.compact_rule.0, rewrite);

    PendingTx::enqueue(scope_be32, rule_id, payload)
    FILE: crates/warp-core/src/scheduler.rs:331-355

CASE 1: Duplicate (scope_hash, rule_id) -
LAST WINS
    index.get(&key) → Some(&i)
    fat[thin[i].handle] = Some(payload) // Overwrite
    thin[i].nonce = next_nonce++      // Refresh nonce

CASE 2: New entry
    fat.push(Some(payload))
    thin.push(RewriteThin { scope_be32, rule_id, nonce, handle })
    index.insert(key, thin.len() - 1)

```

3.3 PendingRewrite Structure

File: crates/warp-core/src/scheduler.rs:68-82

```
pub(crate) struct PendingRewrite {
    pub rule_id: Hash,                                // 32-byte rule identifier
    pub compact_rule: CompactRuleId,                  // u32 hot-path handle
    pub scope_hash: Hash,                            // 32-byte ordering key
    pub scope: NodeKey,                           // { warp_id, local_id }
    pub footprint: Footprint,                      // Read/write declaration
    pub phase: RewritePhase,                     // State machine: Matched → Reserved → ...
}

```

4. Scheduler: Drain & Reserve

4.1 Drain Phase (Radix Sort)

Entry Point: RadixScheduler::drain_for_tx() File: crates/warp-core/src/scheduler.rs:109-110

```
pub(crate) fn drain_for_tx(&mut self, tx: TxId) -> Vec<PendingRewrite> {
    self.pending
        .remove(&tx)
        .map_or_else(Vec::new, |mut txq| txq.drain_in_order())
}
```

Complete Call Trace:

```
RadixScheduler::drain_for_tx(tx)

    self.pending.remove(&tx) -> Option<PendingTx<PendingRewrite>>

PendingTx::drain_in_order()
FILE: crates/warp-core/src/scheduler.rs:416-446

DECISION: n <= 1024 (SMALL_SORT_THRESHOLD)?
YES: sort_unstable_by(cmp_thin)
    Rust std comparison sort

NO: radix_sort()
FILE: crates/warp-core/src/scheduler.rs:360-413

radix_sort()

    Initialize scratch buffer: self.scratch.resize(n, default)

    Lazy allocate histogram: self.counts16 = vec![0u32; 65536]

    FOR pass IN 0..20: // 20 PASSES
```

```

SELECT src/dst buffers (ping-pong)
    flip = false: src=thin, dst=scratch
    flip = true:   src=scratch, dst=thin

PHASE 1: COUNT BUCKETS
FOR r IN src:
    b = bucket16(r, pass)
    counts[b] += 1

PHASE 2: PREFIX SUMS
sum = 0
FOR c IN counts:
    t = *c
    *c = sum
    sum += t

PHASE 3: STABLE SCATTER
FOR r IN src:
    b = bucket16(r, pass)
    dst[counts[b]] = r
    counts[b] += 1

flip = !flip

BUCKET EXTRACTION (bucket16):
FILE: crates/warp-core/src/scheduler.rs:481-498

Pass 0: u16_from_u32_le(r.nonce, 0)      // Nonce bytes [0:2]
Pass 1: u16_from_u32_le(r.nonce, 1)      // Nonce bytes [2:4]
Pass 2: u16_from_u32_le(r.rule_id, 0)    // Rule ID bytes [0:2]
Pass 3: u16_from_u32_le(r.rule_id, 1)    // Rule ID bytes [2:4]
Pass 4: u16_be_from_pair32(scope, 15)   // Scope bytes [30:32]
Pass 5: u16_be_from_pair32(scope, 14)   // Scope bytes [28:30]
...
Pass 19: u16_be_from_pair32(scope, 0)    // Scope bytes [0:2] (MSD)

SORT ORDER: (scope_hash, rule_id, nonce) ascending lexicographic

```

4.2 Reserve Phase (Independence Check)

Entry Point: RadixScheduler::reserve() File: crates/warp-core/src/scheduler.rs:134-143

```

pub(crate) fn reserve(&mut self, tx: TxId, pr: &mut PendingRewrite) -> bool {
    let active = self.active.entry(tx).or_insert_with(ActiveFootprints::new);
    if Self::has_conflict(active, pr) {
        return Self::on_conflict(pr);
    }
}

```

```

    }
    Self::mark_all(active, pr);
    Self::on_reserved(pr)
}

```

Complete Call Trace:

```

RadixScheduler::reserve(tx, pr)

self.active.entry(tx).or_insert_with(ActiveFootprints::new)
    TYPE: HashMap<TxId, ActiveFootprints>
    ActiveFootprints contains 7 GenSets:
        - nodes_written: GenSet<NodeKey>
        - nodes_read: GenSet<NodeKey>
        - edges_written: GenSet<EdgeKey>
        - edges_read: GenSet<EdgeKey>
        - attachments_written: GenSet<AttachmentKey>
        - attachments_read: GenSet<AttachmentKey>
        - ports: GenSet<PortKey>

has_conflict(active, pr) → bool
FILE: crates/warp-core/src/scheduler.rs:157-236

FOR node IN pr.footprint.n_write:
    IF active.nodes_written.contains(node): return true // W-W conflict
    IF active.nodes_read.contains(node): return true // W-R conflict

FOR node IN pr.footprint.n_read:
    IF active.nodes_written.contains(node): return true // R-W conflict
    (R-R is allowed)

FOR edge IN pr.footprint.e_write:
    IF active.edges_written.contains(edge): return true
    IF active.edges_read.contains(edge): return true

FOR edge IN pr.footprint.e_read:
    IF active.edges_written.contains(edge): return true

FOR key IN pr.footprint.a_write:
    IF active.attachments_written.contains(key): return true
    IF active.attachments_read.contains(key): return true

FOR key IN pr.footprint.a_read:
    IF active.attachments_written.contains(key): return true

FOR port IN pr.footprint.b_in  pr.footprint.b_out:
    IF active.ports.contains(port): return true

```

```

IF conflict:
    on_conflict(pr)
        FILE: crates/warp-core/src/scheduler.rs:145-149
        pr.phase = RewritePhase::Aborted
        return false

mark_all(active, pr)
    FILE: crates/warp-core/src/scheduler.rs:238-278

    FOR node IN pr.footprint.n_write:
        active.nodes_written.mark(NodeKey { warp_id, local_id: node })

    FOR node IN pr.footprint.n_read:
        active.nodes_read.mark(NodeKey { ... })

    FOR edge IN pr.footprint.e_write:
        active.edges_written.mark(EdgeKey { ... })

    FOR edge IN pr.footprint.e_read:
        active.edges_read.mark(EdgeKey { ... })

    FOR key IN pr.footprint.a_write:
        active.attachments_written.mark(key)

    FOR key IN pr.footprint.a_read:
        active.attachments_read.mark(key)

    FOR port IN pr.footprint.b_in   pr.footprint.b_out:
        active.ports.mark(port)

on_reserved(pr)
    FILE: crates/warp-core/src/scheduler.rs:151-155
    pr.phase = RewritePhase::Reserved
    return true

```

4.3 GenSet: O(1) Conflict Detection

File: crates/warp-core/src/scheduler.rs:509-535

```

pub(crate) struct GenSet<K> {
    gen: u32,                                // Current generation
    seen: FxHashMap<K, u32>,                  // Key → generation when marked
}

impl<K: Hash + Eq + Copy> GenSet<K> {

```

```

#[inline]
pub fn contains(&self, key: K) -> bool {
    matches!(self.seen.get(&key), Some(&g) if g == self.gen)
}

#[inline]
pub fn mark(&mut self, key: K) {
    self.seen.insert(key, self.gen);
}
}

```

Key Insight: No clearing needed between transactions. Increment `gen` → all old entries become stale.

5. BOAW Parallel Execution

Entry Point: `execute_parallel()` File: `crates/warp-core/src/boaw/exec.rs:61-83`

5.1 Entry Point

```

pub fn execute_parallel(view: GraphView<'_>, items: &[ExecItem], workers: usize) -> Vec<
    assert!(workers >= 1);
    let capped_workers = workers.min(NUM_SHARDS); // Cap at 256

    #[cfg(feature = "parallel-stride-fallback")]
    if std::env::var("ECHO_PARALLEL_STRIDE").is_ok() {
        return execute_parallel_stride(view, items, capped_workers);
    }

    execute_parallel_sharded(view, items, capped_workers) // DEFAULT
}

```

5.2 Complete Call Trace

```

execute_parallel(view, items, workers)

execute_parallel_sharded(view, items, capped_workers)
FILE: crates/warp-core/src/boaw/exec.rs:101-152

IF items.is_empty():
    return (0..workers).map(|_| TickDelta::new()).collect()

partition_into_shards(items.to_vec()) -> Vec<VirtualShard>
FILE: crates/warp-core/src/boaw/shard.rs:109-120

```

```

Create 256 empty VirtualShard structures

FOR item IN items:

    shard_of(&item.scope) → usize
    FILE: crates/warp-core/src/boaw/shard.rs:82-92
    CODE:
        let bytes = scope.as_bytes();
        let first_8: [u8; 8] = [bytes[0..8]];
        let val = u64::from_le_bytes(first_8);
        (val & 255) as usize // SHARD_MASK = 255

    shards[shard_id].items.push(item)

let next_shard = AtomicUsize::new(0)

std::thread::scope(|s| { ... })
FILE: Rust std (scoped threads)

FOR _ IN 0..workers:

    s.spawn(move || { ... }) // WORKER THREAD

    let mut delta = TickDelta::new()
    FILE: crates/warp-core/src/tick_delta.rs:44-52
    Creates: { ops: Vec::new(), origins: Vec::new() }

LOOP: // Work-stealing loop

    shard_id = next_shard.fetch_add(1, Ordering::Relaxed)
    ATOMIC: Returns old value, increments counter
    ORDERING: Relaxed (no synchronization cost)

    IF shard_id >= 256: break

    FOR item IN &shards[shard_id].items:

        let mut scoped = delta.scoped(item.origin)
        FILE: crates/warp-core/src/tick_delta.rs:140-142
        Creates: ScopedDelta { inner: &mut delta, origin, next_op_ix: 0 }

        (item.exec)(view, &item.scope, scoped.inner_mut())

        INSIDE EXECUTOR:
            scoped.emit(op)

```

```

FILE: crates/warp-core/src/tick_delta.rs:234-239
CODE:
    origin.op_ix = self.next_op_ix;
    self.next_op_ix += 1;
    self.inner.emit_with_origin(op, origin);

    TickDelta::emit_with_origin(op, origin)
FILE: crates/warp-core/src/tick_delta.rs:69-75
CODE:
    self.ops.push(op);
    self.origins.push(origin); // if delta_validate

COLLECT THREADS:
handles.into_iter().map(|h| h.join()).collect()
RETURNS: Vec<TickDelta> (one per worker)

```

5.3 ExecItem Structure

File: crates/warp-core/src/boaw/exec.rs:19-35

```

#[derive(Clone, Copy)]
pub struct ExecItem {
    pub exec: ExecuteFn,      // fn(GraphView, &NodeId, &mut TickDelta)
    pub scope: NodeId,        // 32-byte node identifier
    pub origin: OpOrigin,     // { intent_id, rule_id, match_ix, op_ix }
}

```

5.4 Thread Safety

Type	Safety	Reason
GraphView	Sync + Send + Clone	Read-only snapshot
ExecItem	Sync + Send + Copy	Function pointer + primitives
TickDelta	Per-worker exclusive	No shared mutation
AtomicUsize	Lock-free	fetch_add with Relaxed ordering

6. Delta Merge & State Finalization

6.1 Canonical Merge

Entry Point: `merge_deltas()` File: crates/warp-core/src/boaw/merge.rs:36-75
`merge_deltas(deltas: Vec<TickDelta>) → Result<Vec<WarpOp>, MergeConflict>`

```
[1] FLATTEN ALL OPS WITH ORIGINS
    let mut flat: Vec<(WarpOpKey, OpOrigin, WarpOp)> = Vec::new();
    FOR d IN deltas:
        let (ops, origins) = d.into_parts_unsorted();
        FOR (op, origin) IN ops.zip(origins):
            flat.push((op.sort_key(), origin, op));

[2] CANONICAL SORT
    flat.sort_by(|a, b| (&a.0, &a.1).cmp(&(&b.0, &b.1)));
    ORDER: (WarpOpKey, OpOrigin) lexicographic

[3] DEDUPE & CONFLICT DETECTION
    let mut out = Vec::new();
    let mut i = 0;
    WHILE i < flat.len():

        GROUP by WarpOpKey
        key = flat[i].0
        start = i
        WHILE i < flat.len() && flat[i].0 == key: i++

        CHECK if all ops identical
        first = &flat[start].2
        all_same = flat[start+1..i].iter().all(|(_, _, op)| op == first)

        IF all_same:
            out.push(first.clone()) // Accept one copy
        ELSE:
            writers = flat[start..i].iter().map(|(_, o, _)| *o).collect()
            return Err(MergeConflict { writers }) // CONFLICT!

    return Ok(out)
```

6.2 WarpOp Sort Key

File: crates/warp-core/src/tick_patch.rs:207-287

```
pub(crate) fn sort_key(&self) -> WarpOpKey {
    match self {
        Self::OpenPortal { .. } => WarpOpKey { kind: 1, ... },
        Self::UpsertWarpInstance { .. } => WarpOpKey { kind: 2, ... },
        Self::DeleteWarpInstance { .. } => WarpOpKey { kind: 3, ... },
        Self::DeleteEdge { .. } => WarpOpKey { kind: 4, ... }, // Delete before upsert
        Self::DeleteNode { .. } => WarpOpKey { kind: 5, ... },
        Self::UpsertNode { .. } => WarpOpKey { kind: 6, ... },
        Self::UpsertEdge { .. } => WarpOpKey { kind: 7, ... },
```

```

    Self::SetAttachment { ... }      => WarpOpKey { kind: 8, ... }, // Last
}
}

```

Canonical Order: 1. OpenPortal (creates child instances) 2. UpsertWarpInstance 3. DeleteWarpInstance 4. DeleteEdge (delete before upsert) 5. DeleteNode (delete before upsert) 6. UpsertNode 7. UpsertEdge 8. SetAttachment (after skeleton exists)

6.3 State Mutation Methods

File: crates/warp-core/src/graph.rs

GraphStore::insert_node(id, record)
LINE: 175-177
CODE: self.nodes.insert(id, record)

GraphStore::upsert_edge_record(from, edge)
LINE: 196-261
UPDATES:

- self.edge_index.insert(edge_id, from)
- self.edge_to_index.insert(edge_id, to)
- Remove old edge from previous bucket if exists
- self.edges_from.entry(from).or_default().push(edge)
- self.edges_to.entry(to).or_default().push(edge_id)

GraphStore::delete_node_cascade(node)
LINE: 277-354
CASCADES:

- Remove from self.nodes
- Remove node attachment
- Remove ALL outbound edges (and their attachments)
- Remove ALL inbound edges (and their attachments)
- Maintain all 4 index maps consistently

GraphStore::delete_edge_exact(from, edge_id)
LINE: 360-412
VALIDATES: edge is in correct "from" bucket
REMOVES:

- From edges_from bucket
- From edge_index
- From edge_to_index
- From edges_to bucket
- Edge attachment

GraphStore::set_node_attachment(id, value)
LINE: 125-134

```

CODE:
None → self.node_attachments.remove(&id)
Some(v) → self.node_attachments.insert(id, v)

GraphStore::set_edge_attachment(id, value)
LINE: 163-172
Same pattern as node attachments

```

7. Hash Computation

7.1 State Root

Entry Point: compute_state_root() **File:** crates/warp-core/src/snapshot.rs:88-209

```
compute_state_root(state: &WarpState, root: &NodeKey) → Hash
```

[1] BFS REACHABILITY TRAVERSAL

```

Initialize:
reachable_nodes: BTreeSet<NodeKey> = { root }
reachable_warps: BTreeSet<WarpId> = { root.warp_id }
queue: VecDeque<NodeKey> = [ root ]

WHILE let Some(current) = queue.pop_front():

    store = state.store(&current.warp_id)

    FOR edge IN store.edges_from(&current.local_id):
        to = NodeKey { warp_id: current.warp_id, local_id: edge.to }
        IF reachable_nodes.insert(to): queue.push_back(to)

        IF edge has Descend(child_warp) attachment:
            enqueue_descend(state, child_warp, ...)
            Adds child instance root to queue

        IF current node has Descend(child_warp) attachment:
            enqueue_descend(state, child_warp, ...)
```

[2] HASHING PHASE

```

let mut hasher = Hasher::new() // BLAKE3

HASH ROOT BINDING:
hasher.update(&root.warp_id.0) // 32 bytes
```

```

    hasher.update(&root.local_id.0)      // 32 bytes

FOR warp_id IN reachable_warps: // BTreeset = sorted order

    HASH INSTANCE HEADER:
        hasher.update(&instance.warp_id.0)      // 32 bytes
        hasher.update(&instance.root_node.0)    // 32 bytes
        hash_attachment_key_opt(&mut hasher, instance.parent.as_ref())

FOR (node_id, node) IN store.nodes: // BTreeset = sorted
    IF reachable_nodes.contains(&NodeKey { warp_id, local_id: node_id }):
        hasher.update(&node_id.0)            // 32 bytes
        hasher.update(&node.ty.0)           // 32 bytes
        hash_attachment_value_opt(&mut hasher, store.node_attachment(node_id))

FOR (from, edges) IN store.edges_from: // BTreeset = sorted
    IF from is reachable:
        sorted_edges = edges.filter(reachable).sort_by(|a,b| a.id.cmp(b.id))
        hasher.update(&from.0)              // 32 bytes
        hasher.update(&(sorted_edges.len() as u64).to_le_bytes()) // 8 bytes
        FOR edge IN sorted_edges:
            hasher.update(&edge.id.0)        // 32 bytes
            hasher.update(&edge.ty.0)        // 32 bytes
            hasher.update(&edge.to.0)        // 32 bytes
            hash_attachment_value_opt(&mut hasher, store.edge_attachment(&edge.id))

hasher.finalize().into() // → [u8; 32]

```

7.2 Commit Hash v2

Entry Point: `compute_commit_hash_v2()` File: `crates/warp-core/src/snapshot.rs:244-263`

```

pub(crate) fn compute_commit_hash_v2(
    state_root: &Hash,
    parents: &[Hash],
    patch_digest: &Hash,
    policy_id: u32,
) -> Hash {
    let mut h = Hasher::new();
    h.update(&2u16.to_le_bytes());           // Version tag (2 bytes)
    h.update(&(parents.len() as u64).to_le_bytes()); // Parent count (8 bytes)
    for p in parents {
        h.update(p);                      // Each parent (32 bytes)
    }
    h.update(state_root);                 // Graph hash (32 bytes)
    h.update(patch_digest);             // Ops hash (32 bytes)
}

```

```

    h.update(&policy_id.to_le_bytes());           // Policy (4 bytes)
    h.finalize().into()
}

```

Byte Layout:

Offset	Size	Field
0	2	version_tag (0x02 0x00)
2	8	parent_count (u64 LE)
10	32*N	parents[] (N parent hashes)
10+32N	32	state_root
42+32N	32	patch_digest
74+32N	4	policy_id (u32 LE)

TOTAL: 78 + 32*N bytes → BLAKE3 → 32-byte hash

7.3 Patch Digest

Entry Point: `compute_patch_digest_v2()` File: `crates/warp-core/src/tick_patch.rs:755-774`

```

fn compute_patch_digest_v2(
    policy_id: u32,
    rule_pack_id: &ContentHash,
    commit_status: TickCommitStatus,
    in_slots: &[SlotId],
    out_slots: &[SlotId],
    ops: &[WarpOp],
) -> ContentHash {
    let mut h = Hasher::new();
    h.update(&2u16.to_le_bytes());           // Format version
    h.update(&policy_id.to_le_bytes());      // 4 bytes
    h.update(rule_pack_id);                // 32 bytes
    h.update(&[commit_status.code()]);      // 1 byte
    encode_slots(&mut h, in_slots);
    encode_slots(&mut h, out_slots);
    encode_ops(&mut h, ops);
    h.finalize().into()
}

```

8. Commit Orchestration

Entry Point: `Engine::commit_with_receipt()` File: `crates/warp-core/src/engine_impl.rs:837-954`

8.1 Complete Call Trace

```

Engine::commit_with_receipt(tx) → Result<(Snapshot, TickReceipt, WarpTickPatchV1), Eng
[1] VALIDATE TRANSACTION
    IF tx.value() == 0 || !self.live_txs.contains(&tx.value()):
        return Err(EngineError::UnknownTx)

[2] DRAIN CANDIDATES
    policy_id = self.policy_id                                // Line 844
    rule_pack_id = self.compute_rule_pack_id()                // Line 845

    compute_rule_pack_id()
    FILE: engine_impl.rs:1675-1688
    CODE:
        ids = self.rules.values().map(|r| r.id).collect()
        ids.sort_unstable(); ids.dedup()
        hasher.update(&1u16.to_le_bytes()) // version
        hasher.update(&(ids.len() as u64).to_le_bytes())
        FOR id IN ids: hasher.update(&id)
        hasher.finalize().into()

    drained = self.scheduler.drain_for_tx(tx)                  // Line 847
    plan_digest = compute_plan_digest(&drained)              // Line 848

[3] RESERVE (INDEPENDENCE CHECK)
    ReserveOutcome { receipt, reserved, in_slots, out_slots }
    = self.reserve_for_receipt(tx, drained)                  // Line 850-855

    reserve_for_receipt(tx, drained)
    FILE: engine_impl.rs:970-1042

    FOR rewrite IN drained (canonical order):
        accepted = self.scheduler.reserve(tx, &mut rewrite)

        IF !accepted:
            blockers = find_blocking_rewrites(reserved, &rewrite)

            receipt_entries.push(TickReceiptEntry { ... })

        IF accepted:
            reserved.push(rewrite)
            extend_slots_from_footprint(&mut in_slots, &mut out_slots, ...)

    return ReserveOutcome { receipt, reserved, in_slots, out_slots }

```

```

    rewrites_digest = compute_rewrites_digest(&reserved_rewrites) // Line 858

[4] EXECUTE (PHASE 5 BOAW)
    state_before = self.state.clone()                      // Line 862
    delta_ops = self.apply_reserved_rewrites(reserved, &state_before)?
        apply_reserved_rewrites(rewrites, state_before)
        FILE: engine_impl.rs:1044-1105

    let mut delta = TickDelta::new()

    FOR rewrite IN rewrites:
        executor = self.rule_by_compact(rewrite.compact_rule).executor
        view = GraphView::new(self.state.store(&rewrite.scope.warp_id))
        (executor)(view, &rewrite.scope.local_id, &mut delta)

    let ops = delta.finalize() // Canonical sort

    patch = WarpTickPatchV1::new(policy_id, rule_pack_id, ..., ops)
    patch.apply_to_state(&mut self.state)?
        [delta_validate]: assert_delta_matches_diff(&ops, &diff_ops)

[5] MATERIALIZE
    mat_report = self.bus.finalize()                      // Line 884
    self.last_materialization = mat_report.channels
    self.last_materialization_errors = mat_report.errors

[6] COMPUTE DELTA PATCH
    ops = diff_state(&state_before, &self.state)          // Line 889

    diff_state(before, after)
    FILE: tick_patch.rs:979-1069
    - Canonicalize portal authoring (OpenPortal)
    - Diff instances (delete/upsert)
    - Diff nodes, edges, attachments
    - Sort by WarpOp::sort_key()

    patch = WarpTickPatchV1::new(policy_id, rule_pack_id, ..., ops)
    patch_digest = patch.digest()                         // Line 898

[7] COMPUTE STATE ROOT
    state_root = compute_state_root(&self.state, &self.current_root) // Line 900

[8] GET PARENTS

```

```

parents = self.last_snapshot.as_ref().map(|s| vec![s.hash]).unwrap_or_default()

[9] COMPUTE DECISION DIGEST
    decision_digest = receipt.digest()                                // Line 929

[10] COMPUTE COMMIT HASH
    hash = compute_commit_hash_v2(&state_root, &parents, &patch_digest, policy_id)

[11] BUILD SNAPSHOT
    snapshot = Snapshot {
        root: self.current_root,
        hash,                                         // commit_id v2
        parents,
        plan_digest,                                // Diagnostic
        decision_digest,                            // Diagnostic
        rewrites_digest,                           // Diagnostic
        patch_digest,                               // COMMITTED
        policy_id,                                 // COMMITTED
        tx,
    }

[12] RECORD TO HISTORY
    self.last_snapshot = Some(snapshot.clone())           // Line 947
    self.tick_history.push((snapshot, receipt, patch))   // Line 948-949
    self.live_txs.remove(&tx.value())                   // Line 951
    self.scheduler.finalize_tx(tx)                      // Line 952

[13] RETURN
    Ok((snapshot, receipt, patch))

```

8.2 Commit Hash Inputs

Input	Committed?	Purpose
state_root		What the graph looks like
patch_digest		How we got here (ops)
parents		Chain continuity
policy_id		Aion policy version
plan_digest		Diagnostic only
decision_digest		Diagnostic only
rewrites_digest		Diagnostic only

9. Complete Call Graph

9.1 Full Journey: Intent → Commit

USER ACTION

```

Engine::ingest_intent(intent_bytes)
    compute_intent_id()                                // BLAKE3 content hash
    make_node_id(), make_type_id()                    // Structural IDs
    store.insert_node()                               // Create event node
    store.set_node_attachment()                      // Attach intent payload
    store.insert_edge()                               // Pending edge to inbox

Engine::begin() → TxId
    tx_counter.wrapping_add(1)
    live_txs.insert(tx_counter)
    TxId::from_raw(tx_counter)

Engine::dispatch_next_intent(tx)                   // (or manual apply)

Engine::apply(tx, rule_name, scope)
    Engine::apply_in_warp(tx, warp_id, rule_name, scope, &[])
        rules.get(rule_name)                         // Lookup rule
        GraphView::new(store)                        // Read-only view
        (rule.matcher)(view, scope)                  // Match check
        scope_hash()                                // BLAKE3 ordering key
        (rule.compute_footprint)(view, scope)        // Footprint
        scheduler.enqueue(tx, PendingRewrite)
            PendingTx::enqueue()                     // Last-wins dedup

Engine::commit_with_receipt(tx)

[DRAIN]
    scheduler.drain_for_tx(tx)
        PendingTx::drain_in_order()
            radix_sort() or sort_unstable_by()
                20-pass LSD radix sort
                ORDER: (scope_hash, rule_id, nonce)

[RESERVE]
    FOR rewrite IN drained:

```

```

scheduler.reserve(tx, &mut rewrite)
    has_conflict(active, pr)
        GenSet::contains() × N      // O(1) per check
    mark_all(active, pr)
        GenSet::mark() × M          // O(1) per mark

[EXECUTE]
apply_reserved_rewrites(reserved, state_before)
    FOR rewrite IN reserved:
        (executor)(view, &scope, &mut delta)
            scoped.emit(op)
                delta.emit_with_origin(op, origin)
            delta.finalize()           // Sort ops
        patch.apply_to_state(&mut self.state)

[MATERIALIZE]
bus.finalize()

[DELTA PATCH]
diff_state(&state_before, &self.state)
    Sort by WarpOp::sort_key()
WarpTickPatchV1::new(...)
    compute_patch_digest_v2()

[HASHES]
compute_state_root(&self.state, &self.current_root)
    BFS reachability
    BLAKE3 over canonical encoding
compute_commit_hash_v2(state_root, parents, patch_digest, policy_id)
    BLAKE3(version || parents || state_root || patch_digest || policy_id)

[SNAPSHOT]
Snapshot { root, hash, parents, digests..., policy_id, tx }

[RECORD]
tick_history.push((snapshot, receipt, patch))
live_txs.remove(&tx.value())
scheduler.finalize_tx(tx)

RETURN: (Snapshot, TickReceipt, WarpTickPatchV1)

```

9.2 File Index

Component	Primary File	Key Lines
Intent Ingestion	<code>engine_impl.rs</code>	1216-1281
Identity Hashing	<code>ident.rs</code>	85-109
Transaction Begin	<code>engine_impl.rs</code>	711-719
Rule Apply	<code>engine_impl.rs</code>	730-806
Footprint	<code>footprint.rs</code>	131-152
Scheduler Enqueue	<code>scheduler.rs</code>	102-105, 331-355
Radix Sort	<code>scheduler.rs</code>	360-413, 481-498
Reserve/Conflict	<code>scheduler.rs</code>	134-278
GenSet	<code>scheduler.rs</code>	509-535
BOAW Execute	<code>boaw/exec.rs</code>	61-152
Shard Routing	<code>boaw/shard.rs</code>	82-120
Delta Merge	<code>boaw/merge.rs</code>	36-75
TickDelta	<code>tick_delta.rs</code>	38-172
WarpOp Sort Key	<code>tick_patch.rs</code>	207-287
State Mutations	<code>graph.rs</code>	175-412
Patch Apply	<code>tick_patch.rs</code>	434-561
Diff State	<code>tick_patch.rs</code>	979-1069
State Root Hash	<code>snapshot.rs</code>	88-209
Commit Hash v2	<code>snapshot.rs</code>	244-263
Patch Digest	<code>tick_patch.rs</code>	755-774
Commit Orchestrator	<code>engine_impl.rs</code>	837-954

Appendix A: Complexity Summary

Operation	Complexity	Notes
<code>ingest_intent</code>	$O(1)$	Fixed structural insertions
<code>begin</code>	$O(1)$	Counter increment + set insert
<code>apply</code>	$O(m)$	$m = \text{footprint size}$
<code>drain_for_tx</code> (radix)	$O(n)$	$n = \text{candidates}, 20 \text{ passes}$
reserve per rewrite	$O(m)$	$m = \text{footprint size}, O(1) \text{ per check}$
<code>execute_parallel</code>	$O(n/w)$	$n = \text{items}, w = \text{workers}$
<code>merge_deltas</code>	$O(k \log k)$	$k = \text{total ops (sort + dedup)}$
<code>compute_state_root</code>	$O(V + E)$	$V = \text{nodes}, E = \text{edges}$
<code>compute_commit_hash_v2</code>	$O(P)$	$P = \text{parents}$

Appendix B: Determinism Boundaries

Guaranteed Deterministic

- Radix sort ordering (20-pass LSD)
- BTreeMap/BTreeSet iteration
- BLAKE3 hashing
- GenSet conflict detection
- Canonical merge deduplication

Intentionally Non-Deterministic (Handled by Merge)

- Worker execution order in BOAW
- Shard claim order (atomic counter)

Protocol Constants (Frozen)

- NUM_SHARDS = 256
- SHARD_MASK = 255
- Shard routing: `LE_u64(node_id[0..8]) & 255`
- Commit hash v2 version tag: 0x02 0x00

Document generated 2026-01-18. File paths and line numbers accurate as of this date.